

FAISS

🕒 Created	@February 14, 2025 12:51 PM
🏷️ Tags	Personal

PDF Text Processing and Embedding Storage with FAISS

This code demonstrates how to extract text from a PDF, split it into chunks, generate embeddings for each chunk using Sentence Transformers, and then store these embeddings in a FAISS index for efficient similarity search.

1. PDF Text Extraction

```
```python
import fitz # PyMuPDF

def extract_text_from_pdf(pdf_path):
 doc = fitz.open(pdf_path)
 text = "\n".join([page.get_text() for page in doc])
 return text

pdf_text = extract_text_from_pdf("/Users/vinod/Desktop/mike/sample.pdf")
```

### 2. Text Chunking

```
from langchain.text_splitter import RecursiveCharacterTextSplitter

def split_text(text, chunk_size=500, chunk_overlap=50):
 text_splitter = RecursiveCharacterTextSplitter(
 chunk_size=chunk_size, chunk_overlap=chunk_overlap
)
 return text_splitter.split_text(text)
```

```
chunks = split_text(pdf_text)
print(f"Total chunks: {len(chunks)}")
```

### 3. Embedding Generation

```
from sentence_transformers import SentenceTransformer

model = SentenceTransformer("sentence-transformers/all-MiniLM-L6-v2")
chunk_embeddings = model.encode(chunks)

Print first 5 embeddings (optional)
for i, embedding in enumerate(chunk_embeddings[:5]):
 print(f"Embedding {i+1}: {embedding}\n")

print(f"Generated {len(chunk_embeddings)} embeddings!")
```

- `from sentence_transformers import SentenceTransformer` : Imports the Sentence Transformer library. Install it: `pip install sentence-transformers` .
- `model = ...` : Loads a pre-trained Sentence Transformer model. `all-MiniLM-L6-v2` is a good balance of speed and performance.
- `chunk_embeddings = model.encode(chunks)` : Generates embeddings for each text chunk.
- `for i, embedding in enumerate(chunk_embeddings[:5]):` : (Optional) Prints the first 5 embeddings for inspection. Remove this if you don't need to see the embeddings.
- `print(...)` : Prints the number of embeddings generated.

### 4. FAISS Indexing

```
import faiss
import numpy as np
```

```

dimension = chunk_embeddings.shape[1] # 384 for MiniLM
index = faiss.IndexFlatL2(dimension) # L2 = Euclidean distance

Convert embeddings to FAISS format
faiss_data = np.array(chunk_embeddings, dtype=np.float32)
index.add(faiss_data)

print("Embeddings stored in FAISS index!")

```

- `import faiss` : Imports the FAISS library. Install it: `pip install faiss-cpu` (for CPU) or `pip install faiss-gpu` (if you have a compatible GPU).
- `import numpy as np` : Imports NumPy for numerical operations.
- `dimension = chunk_embeddings.shape[1]` : Gets the dimensionality of the embeddings (384 for `all-MiniLM-L6-v2` ).
- `index = faiss.IndexFlatL2(dimension)` : Creates a FAISS index. `IndexFlatL2` is a simple and efficient index for storing and searching L2-normalized vectors (embeddings).
- `faiss_data = np.array(chunk_embeddings, dtype=np.float32)` : Converts the embeddings to a NumPy array of `float32` data type, which is required by FAISS.
- `index.add(faiss_data)` : Adds the embeddings to the FAISS index.
- `print(...)` : Confirms that the embeddings have been added to the index.